

# Mining Data Streams (Part 2)

Mining of Massive Datasets

Jure Leskovec, Anand Rajaraman, Jeff Ullman

Stanford University

<http://www.mmds.org>



## ■ More algorithms for streams:

- **(1) Filtering a data stream: Bloom filters**
  - Select elements with property  $x$  from stream
- **(2) Counting distinct elements: Flajolet-Martin**
  - Number of distinct elements in the last  $k$  elements of the stream
- **(3) Estimating moments: AMS method**
  - Estimate std. dev. of last  $k$  elements
- **(4) Counting frequent items**

# **(1) Filtering Data Streams**

# Filtering Data Streams

- Each element of data stream is a tuple
- Given a list of keys  $S$
- Determine which tuples of stream are in  $S$
- Obvious solution: Hash table
  - But suppose we do not have enough memory to store all of  $S$  in a hash table
    - E.g., we might be processing millions of filters on the same stream

# Applications

## ■ Example: Email spam filtering

- We know 1 billion “good” email addresses
- If an email comes from one of these, it is **NOT** spam

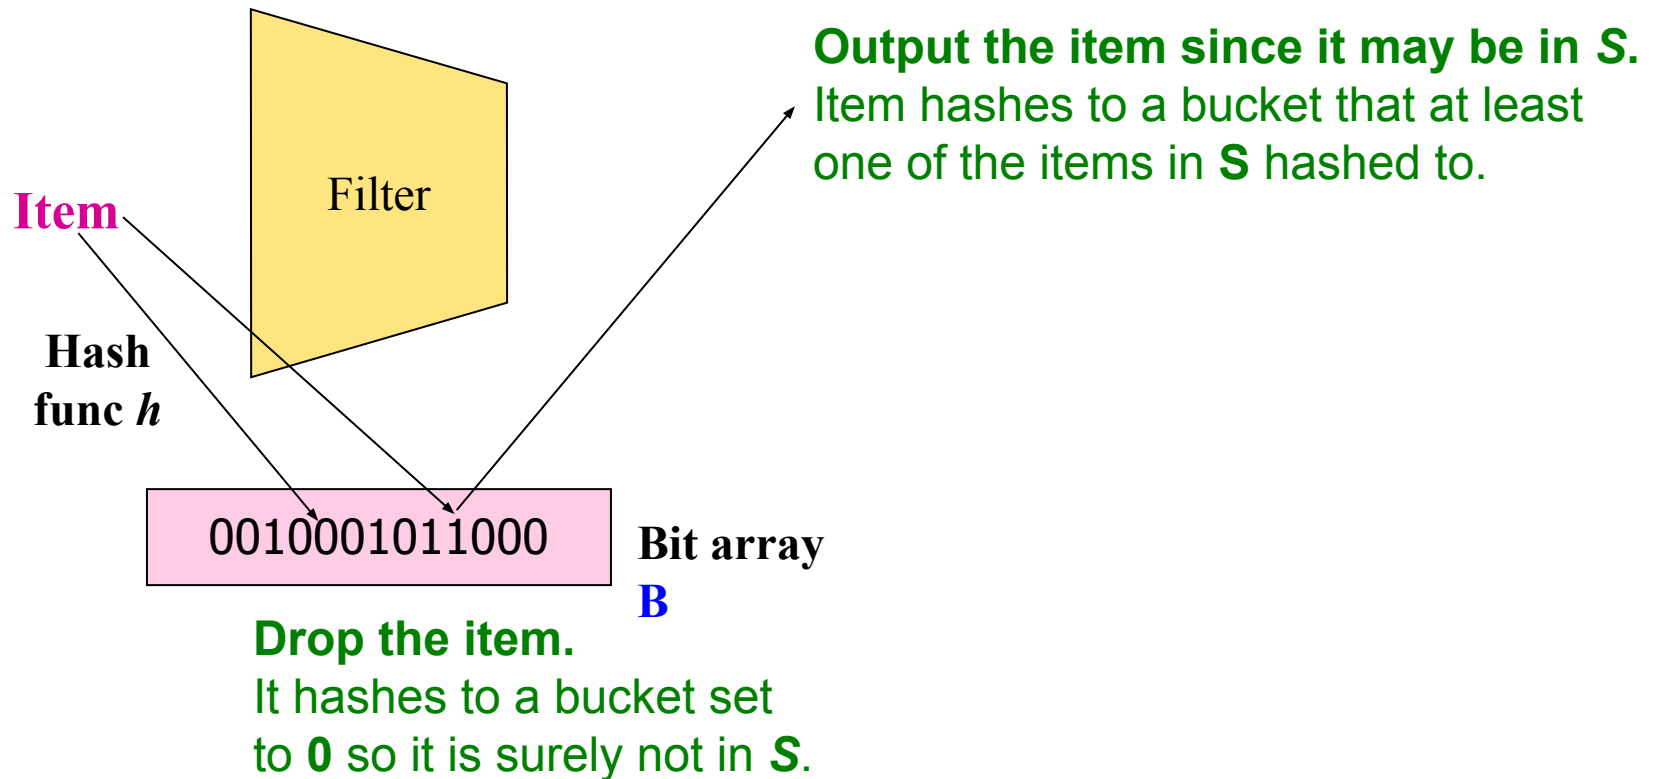
## ■ Publish-subscribe systems

- You are collecting lots of messages (news articles)
- People express interest in certain sets of keywords
- Determine whether each message matches user’s interest

# First Cut Solution (1)

- Given a set of keys  $S$  that we want to filter
- Create a **bit array**  $B$  of  $n$  bits, initially all  $0$ s
- Choose a **hash function**  $h$  with range  $[0, n)$
- Hash each member of  $s \in S$  to one of  $n$  buckets, and set that bit to  $1$ , i.e.,  $B[h(s)] = 1$
- Hash each element  $a$  of the stream and output only those that hash to bit that was set to  $1$ 
  - **Output**  $a$  if  $B[h(a)] == 1$

# First Cut Solution (2)



- **Creates false positives but no false negatives**
  - If the item is in  $S$  we surely output it, if not we may still output it

# First Cut Solution (3)

- $|S| = 1$  billion email addresses  
 $|B| = 1\text{GB} = 8$  billion bits
- If the email address is in  $S$ , then it surely hashes to a bucket that has the bit set to **1**, so it always gets through (*no false negatives*)
- Approximately  $1/8$  of the bits are set to **1**, so about  $1/8^{\text{th}}$  of the addresses not in  $S$  get through to the output (*false positives*)
  - Actually, less than  $1/8^{\text{th}}$ , because more than one address might hash to the same bit




# Analysis: Throwing Darts (1)

- More accurate analysis for the number of false positives
- **Consider:** If we throw  $m$  darts into  $n$  equally likely targets, **what is the probability that a target gets at least one dart?**
- **In our case:**
  - Targets = bits/buckets
  - Darts = hash values of items

# Analysis: Throwing Darts (2)

- We have  $m$  darts,  $n$  targets
- What is the probability that a target gets at least one dart?



# Analysis: Throwing Darts (3)

- **Fraction of 1s in the array B =**  
**= probability of false positive =  $1 - e^{-m/n}$**
- **Example:  $10^9$  darts,  $8 \cdot 10^9$  targets**
  - Fraction of 1s in B =  $1 - e^{-1/8} = 0.1175$ 
    - Compare with our earlier estimate:  $1/8 = 0.125$

# Bloom Filter

- Consider:  $|\mathbf{S}| = m$ ,  $|\mathbf{B}| = n$
- Use  $k$  independent hash functions  $h_1, \dots, h_k$
- Initialization:
  - Set  $\mathbf{B}$  to all 0s
  - Hash each element  $s \in \mathbf{S}$  using each hash function  $h_i$ , set  $\mathbf{B}[h_i(s)] = 1$  (for each  $i = 1, \dots, k$ ) (note: we have a single array B!)
- Run-time:
  - When a stream element with key  $x$  arrives
    - If  $\mathbf{B}[h_i(x)] = 1$  for all  $i = 1, \dots, k$  then declare that  $x$  is in  $\mathbf{S}$ 
      - That is,  $x$  hashes to a bucket set to 1 for every hash function  $h_i(x)$
    - Otherwise discard the element  $x$

# Bloom Filter -- Analysis

- What fraction of the bit vector  $B$  are 1s?
  - Throwing  $k \cdot m$  darts at  $n$  targets
  - So fraction of 1s is  $(1 - e^{-km/n})$
- But we have  $k$  independent hash functions and we only let the element  $x$  through if all  $k$  hash element  $x$  to a bucket of value 1
- So, false positive probability =  $(1 - e^{-km/n})^k$

# Bloom Filter – Analysis (2)

■  $m = 1$  billion,  $n = 8$  billion

■  $k = 1$ :  $(1 - e^{-1/8}) = 0.1175$


■  $k = 2$ :  $(1 - e^{-1/4})^2 = 0.0493$

■ What happens as we keep increasing  $k$ ?

■ “Optimal” value of  $k$ :  $n/m \ln(2)$

■ In our case: Optimal  $k = 8 \ln(2) = 5.54 \approx 6$

■ Error at  $k = 6$ :  $(1 - e^{-1/6})^2 = 0.0235$



# Bloom Filter: Wrap-up

- Bloom filters guarantee no false negatives, and use limited memory
  - Great for pre-processing before more expensive checks
- Suitable for hardware implementation
  - Hash function computations can be parallelized
- Is it better to have 1 big B or  $k$  small Bs?
  - It is the same:  $(1 - e^{-km/n})^k$  vs.  $(1 - e^{-m/(n/k)})^k$
  - But keeping 1 big B is simpler

## **(2) Counting Distinct Elements**



# Counting Distinct Elements

## ■ Problem:

- Data stream consists of a universe of elements chosen from a set of size  $N$
- Maintain a count of the number of distinct elements seen so far

## ■ Obvious approach:

Maintain the set of elements seen so far

- That is, keep a hash table of all the distinct elements seen so far

# Applications

- **How many different words are found among the Web pages being crawled at a site?**
  - Unusually low or high numbers could indicate artificial pages (spam?)
- **How many different Web pages does each customer request in a week?**
- **How many distinct products have we sold in the last week?**

# Using Small Storage

- **Real problem: What if we do not have space to maintain the set of elements seen so far?**
- **Estimate the count in an unbiased way**
- **Accept that the count may have a little error, but limit the probability that the error is large**

# Flajolet-Martin Approach

- Pick a hash function  $h$  that maps each of the  $N$  elements to at least  $\log_2 N$  bits
- For each stream element  $a$ , let  $r(a)$  be the number of trailing 0s in  $h(a)$ 
  - $r(a)$  = position of first 1 counting from the right
    - E.g., say  $h(a) = 12$ , then 12 is 1100 in binary, so  $r(a) = 2$
- Record  $R$  = the maximum  $r(a)$  seen
  - $R = \max_a r(a)$ , over all the items  $a$  seen so far
- Estimated number of distinct elements =  $2^R$

# Why It Works: Intuition

- Very very rough and heuristic intuition why Flajolet-Martin works:
  - $h(a)$  hashes  $a$  with equal prob. to any of  $N$  values
  - Then  $h(a)$  is a sequence of  $\log_2 N$  bits, where  $2^{-r}$  fraction of all  $a$ s have a tail of  $r$  zeros
    - About 50% of  $a$ s hash to  $***0$
    - About 25% of  $a$ s hash to  $**00$
    - So, if we saw the longest tail of  $r=2$  (i.e., item hash ending  $*100$ ) then we have probably seen **about 4** distinct items so far
  - So, it takes to hash about  $2^r$  items before we see one with zero-suffix of length  $r$

# Why It Works: More formally

- Now we show why Flajolet-Martin works
- Formally, we will show that **probability of finding a tail of  $r$  zeros:**
  - Goes to **1** if  $m \gg 2^r$
  - Goes to **0** if  $m \ll 2^r$where  $m$  is the number of distinct elements seen so far in the stream
- **Thus,  $2^R$  will almost always be around  $m$ !**

# Why It Works: More formally

- What is the probability that a given  $h(a)$  ends in at least  $r$  zeros is  $2^{-r}$ 
  - $h(a)$  hashes elements uniformly at random
  - Probability that a random number ends in at least  $r$  zeros is  $2^{-r}$
- Then, the probability of **NOT** seeing a tail of length  $r$  among  $m$  elements:

$$(1 - 2^{-r})^m$$

Prob. all end in fewer than  $r$  zeros.

Prob. that given  $h(a)$  ends in fewer than  $r$  zeros

# Why It Works: More formally

- **Note:**  $(1 - 2^{-r})^m = (1 - 2^{-r})^{2^r (m2^{-r})} \approx e^{-m2^{-r}}$
- **Prob. of NOT finding a tail of length  $r$  is:**
  - If  $m \ll 2^r$ , then prob. tends to **1**
    - $(1 - 2^{-r})^m \approx e^{-m2^{-r}} = 1$  as  $m/2^r \rightarrow 0$
    - So, the probability of finding a tail of length  $r$  tends to **0**
  - If  $m \gg 2^r$ , then prob. tends to **0**
    - $(1 - 2^{-r})^m \approx e^{-m2^{-r}} = 0$  as  $m/2^r \rightarrow \infty$
    - So, the probability of finding a tail of length  $r$  tends to **1**
- **Thus,  $2^R$  will almost always be around  $m$ !**



# Why It Doesn't Work

- $E[2^R]$  is actually infinite
  - Probability halves when  $R \rightarrow R+1$ , but value doubles
- Workaround involves using many hash functions  $h_i$  and getting many samples of  $R_i$
- How are samples  $R_i$  combined?
  - Average? What if one very large value  $2^{R_i}$ ?
  - Median? All estimates are a power of 2
  - Solution:
    - Partition your samples into small groups
    - Take the median of groups
    - Then take the average of the medians

## **(3) Computing Moments**

# Generalization: Moments

- Suppose a stream has elements chosen from a set  $A$  of  $N$  values
- Let  $m_i$  be the number of times value  $i$  occurs in the stream
- The  $k^{\text{th}}$  *moment* is

$$\sum_{i \in A} (m_i)^k$$

# Special Cases

$$\sum_{i \in A} (m_i)^k$$

- **0<sup>th</sup> moment** = number of distinct elements
  - The problem just considered
- **1<sup>st</sup> moment** = count of the numbers of elements = length of the stream
  - Easy to compute
- **2<sup>nd</sup> moment** = *surprise number S* =  
a measure of how uneven the distribution is

# Example: Surprise Number

- Stream of length 100
- 11 distinct values
- Item counts: 10, 9, 9, 9, 9, 9, 9, 9, 9, 9, 9  
Surprise  $S = 910$
- Item counts: 90, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1  
Surprise  $S = 8,110$

# AMS Method

- AMS method works for all moments
- Gives an unbiased estimate
- We will just concentrate on the 2<sup>nd</sup> moment  $S$
- We pick and keep track of many variables  $X$ :
  - For each variable  $X$  we store  $X.el$  and  $X.val$ 
    - $X.el$  corresponds to the item  $i$
    - $X.val$  corresponds to the **count** of item  $i$
  - Note this requires a count in main memory, so number of  $X$ s is limited
- Our goal is to compute  $S = \sum_i m_i^2$

# One Random Variable ( $X$ )

## ■ How to set $X.val$ and $X.el$ ?


- Assume stream has length  $n$  (we relax this later)
- Pick some random time  $t$  ( $t < n$ ) to start, so that any time is equally likely
- Let at time  $t$  the stream have item  $i$ . **We set  $X.el = i$**
- Then we maintain count  $c$  ( **$X.val = c$** ) of the number of  $i$ s in the stream starting from the chosen time  $t$

## ■ Then the estimate of the 2<sup>nd</sup> moment ( $\sum_i m_i^2$ ) is:

$$S = f(X) = n(2 \cdot c - 1)$$

- Note, we will keep track of multiple  $X$ s, ( $X_1, X_2, \dots, X_k$ ) and our final estimate will be  $S = 1/k \sum_j^k f(X_j)$

# Expectation Analysis



■ **2<sup>nd</sup> moment is  $S = \sum_i m_i^2$**

■  **$c_t$**  ... number of times item at time  **$t$**  appears from time  **$t$**  onwards ( $c_1=m_a$ ,  $c_2=m_a-1$ ,  $c_3=m_b$ )

■  $E[f(X)] = \frac{1}{n} \sum_{t=1}^n n(2c_t - 1)$

$= \frac{1}{n} \sum_i n (1 + 3 + 5 + \dots + 2m_i - 1)$

$m_i$  ... total count of item  $i$  in the stream (we are assuming stream has length  $n$ )

Group times by the value seen


Time  $t$  when the last  $i$  is seen ( $c_t=1$ )

Time  $t$  when the penultimate  $i$  is seen ( $c_t=2$ )

Time  $t$  when the first  $i$  is seen ( $c_t=m_i$ )



# Expectation Analysis



- $E[f(X)] = \frac{1}{n} \sum_i n (1 + 3 + 5 + \dots + 2m_i - 1)$ 
  - Little side calculation:  $(1 + 3 + 5 + \dots + 2m_i - 1) = \sum_{i=1}^{m_i} (2i - 1) = 2 \frac{m_i(m_i+1)}{2} - m_i = (m_i)^2$
- Then  $E[f(X)] = \frac{1}{n} \sum_i n (m_i)^2$
- So,  $E[f(X)] = \sum_i (m_i)^2 = S$
- We have the second moment (in expectation)!

# Higher-Order Moments

- For estimating  $k^{\text{th}}$  moment we essentially use the same algorithm but change the estimate:
  - For  $k=2$  we used  $n (2 \cdot c - 1)$
  - For  $k=3$  we use:  $n (3 \cdot c^2 - 3c + 1)$  (where  $c=X.\text{val}$ )
- Why?
  - For  $k=2$ : Remember we had  $(1 + 3 + 5 + \dots + 2m_i - 1)$  and we showed terms  $2c-1$  (for  $c=1, \dots, m$ ) sum to  $m^2$ 
    - $\sum_{c=1}^m 2c - 1 = \sum_{c=1}^m c^2 - \sum_{c=1}^m (c - 1)^2 = m^2$
    - So:  $2c - 1 = c^2 - (c - 1)^2$
  - For  $k=3$ :  $c^3 - (c-1)^3 = 3c^2 - 3c + 1$
- Generally: Estimate =  $n (c^k - (c - 1)^k)$

# Combining Samples

## ■ In practice:

- Compute  $f(X) = n(2c - 1)$  for as many variables  $X$  as you can fit in memory
- Average them in groups
- Take median of averages

## ■ Problem: Streams never end

- We assumed there was a number  $n$ , the number of positions in the stream
- But real streams go on forever, so  $n$  is a variable – the number of inputs seen so far

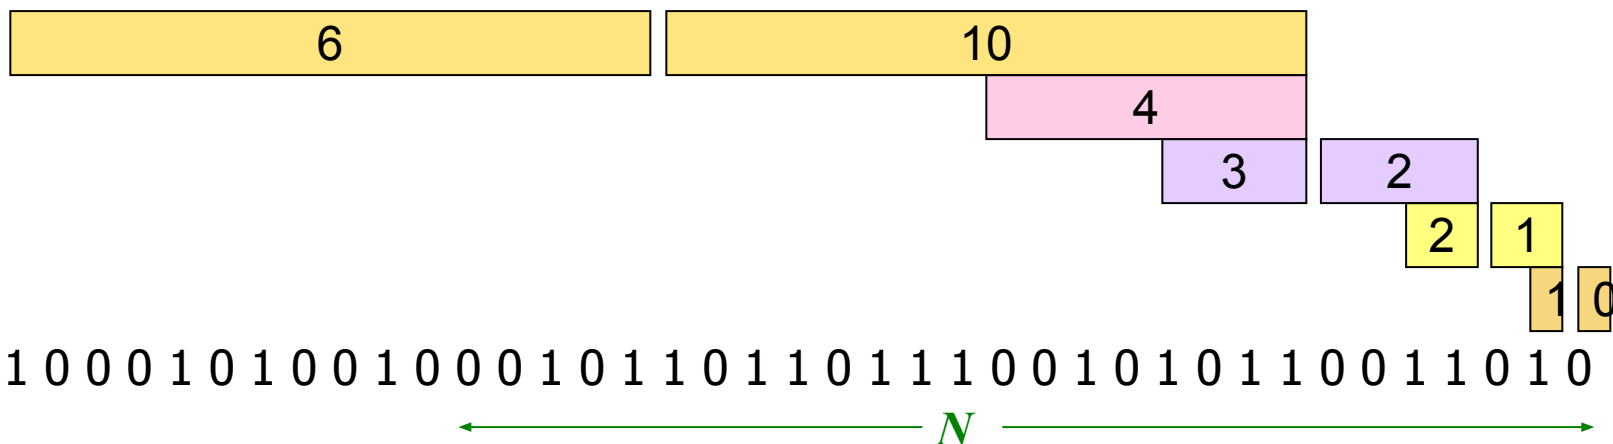
# Streams Never End: Fixups

- **(1)** The variables  $X$  have  $n$  as a factor – keep  $n$  separately; just hold the count in  $X$
- **(2)** Suppose we can only store  $k$  counts. We must throw some  $X$ s out as time goes on:
  - **Objective:** Each starting time  $t$  is selected with probability  $k/n$
  - **Solution: (fixed-size sampling!)**
    - Choose the first  $k$  times for  $k$  variables
    - When the  $n^{\text{th}}$  element arrives ( $n > k$ ), choose it with probability  $k/n$
    - If you choose it, throw one of the previously stored variables  $X$  out, with equal probability

# Counting Itemsets

# Counting Itemsets

- New Problem: Given a stream, which items appear more than  $s$  times in the window?
- **Possible solution**: Think of the stream of baskets as one binary stream per item
  - **1** = item present; **0** = not present
  - Use **DGIM** to estimate counts of **1s** for all items



# Extensions

- In principle, you could count frequent pairs or even larger sets the same way
  - One stream per itemset
- Drawbacks:
  - Only approximate
  - Number of itemsets is way too big

# Exponentially Decaying Windows


- **Exponentially decaying windows: A heuristic for selecting likely frequent item(sets)**
  - **What are “currently” most popular movies?**
    - Instead of computing the raw count in last  $N$  elements
    - Compute a **smooth aggregation** over the whole stream
- If stream is  $a_1, a_2, \dots$  and we are taking the sum of the stream, take the answer at time  $t$  to be:  
$$= \sum_{i=1}^t a_i (1 - c)^{t-i}$$
  - $c$  is a constant, presumably tiny, like  $10^{-6}$  or  $10^{-9}$
- **When new  $a_{t+1}$  arrives:**  
Multiply current sum by  $(1-c)$  and add  $a_{t+1}$



# Example: Counting Items

- If each  $a_i$  is an “item” we can compute the **characteristic function** of each possible item  $x$  as an Exponentially Decaying Window
  - That is:  $\sum_{i=1}^t \delta_i \cdot (1 - c)^{t-i}$   
where  $\delta_i=1$  if  $a_i=x$ , and  $0$  otherwise
  - Imagine that for each item  $x$  we have a binary stream ( $1$  if  $x$  appears,  $0$  if  $x$  does not appear)
  - **New item  $x$  arrives:**
    - Multiply all counts by  $(1-c)$
    - Add  $+1$  to count for element  $x$
- **Call this sum the “weight” of item  $x$**

# Sliding Versus Decaying Windows



- **Important property:** Sum over all weights  $\sum_t (1 - c)^t$  is  $1/[1 - (1 - c)] = 1/c$

# Example: Counting Items

- What are “currently” most popular movies?
- Suppose we want to find movies of weight  $> \frac{1}{2}$ 
  - **Important property:** Sum over all weights  $\sum_t (1 - c)^t$  is  $1/[1 - (1 - c)] = 1/c$
- **Thus:**
  - There cannot be more than  $2/c$  movies with weight of  $\frac{1}{2}$  or more
- So,  $2/c$  is a limit on the number of movies being counted at any time

# Extension to Itemsets

- **Count (some) itemsets in an E.D.W.**
  - **What are currently “hot” itemsets?**
    - **Problem:** Too many itemsets to keep counts of all of them in memory
- **When a basket B comes in:**
  - Multiply all counts by **(1-c)**
  - For uncounted items in **B**, create new count
  - Add **1** to count of any item in **B** and to any **itemset** contained in **B** that is already being counted
  - **Drop counts  $< \frac{1}{2}$**
  - Initiate new counts (next slide)

# Initiation of New Counts

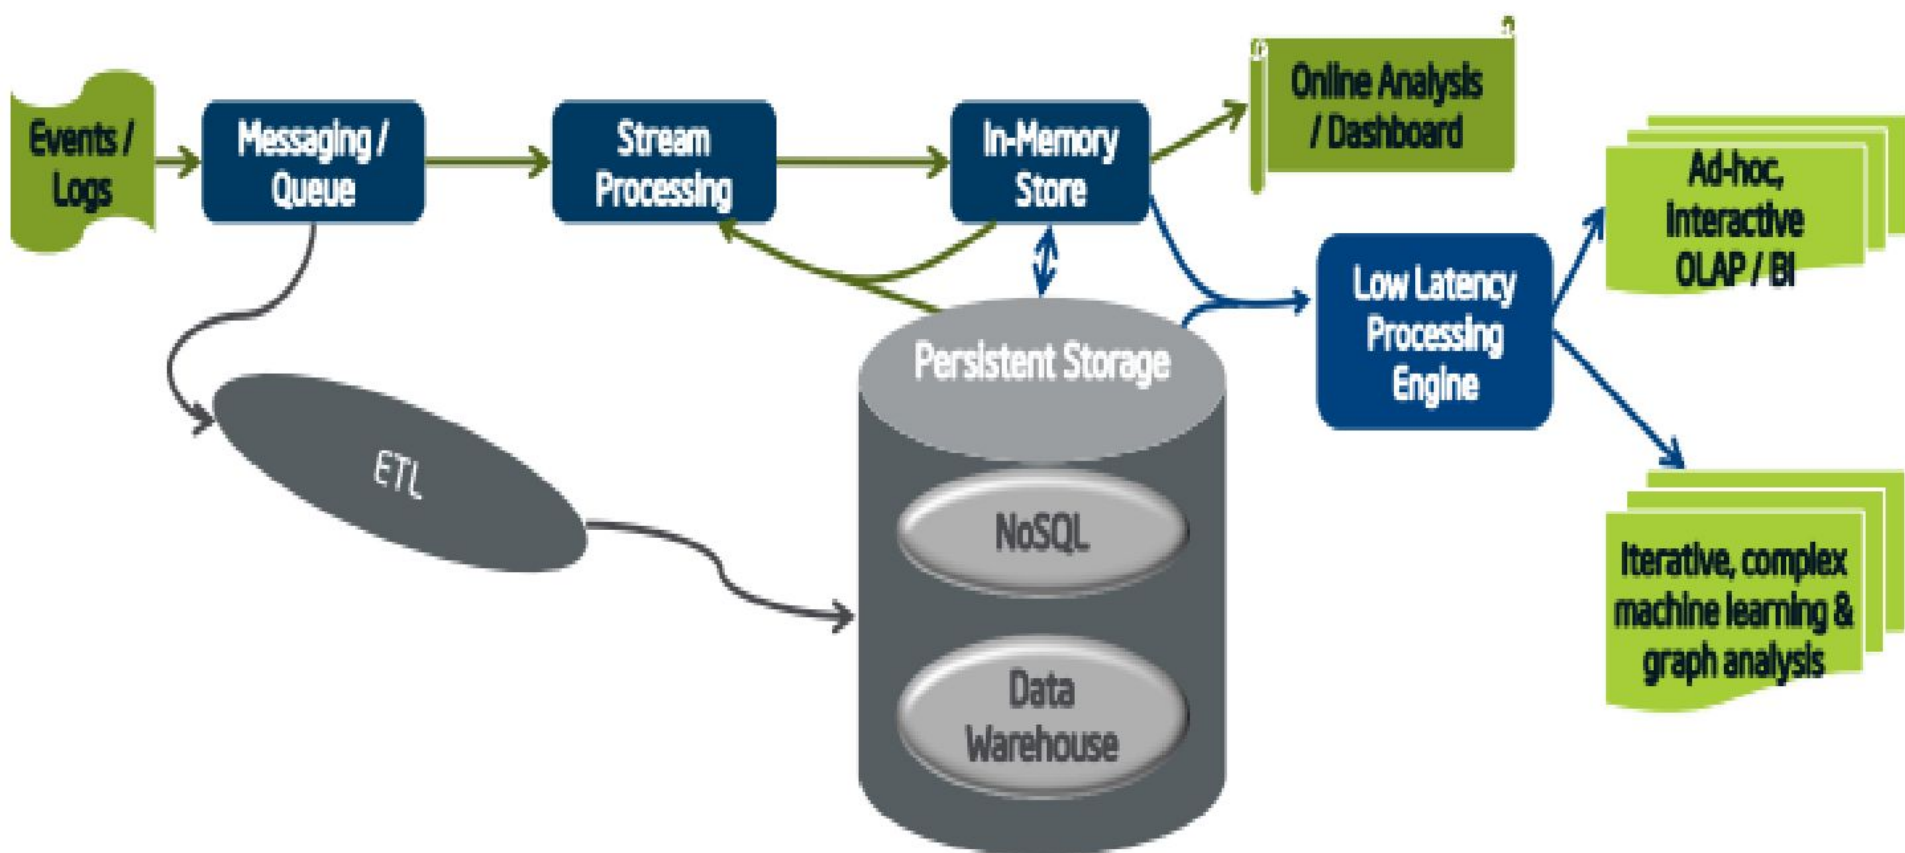
- Start a count for an itemset  $S \subseteq B$  if every proper subset of  $S$  had a count prior to arrival of basket  $B$ 
  - **Intuitively:** If all subsets of  $S$  are being counted this means they are “frequent/hot” and thus  $S$  has a potential to be “hot”
- **Example:**
  - Start counting  $S=\{i, j\}$  iff both  $i$  and  $j$  were counted prior to seeing  $B$
  - Start counting  $S=\{i, j, k\}$  iff  $\{i, j\}$ ,  $\{i, k\}$ , and  $\{j, k\}$  were all counted prior to seeing  $B$

# How many counts do we need?

- Counts for single items  $< (2/c) \cdot (\text{avg. number of items in a basket})$
- Counts for larger itemsets = ??
- But we are conservative about starting counts of large sets
  - If we counted every set we saw, one basket of **20** items would initiate **1M** counts

# Real time Analytics Platform (RTAP) Applications

- A vision for next-gen big data analytics
- Data captured & processed in a (semi) streaming/online fashion
- Real-time & history data combined and mined interactively and/or iteratively – Complex OLAP / BI in interactive fashion – Iterative, complex machine learning & graph analysis
- Predominantly memory-based computation



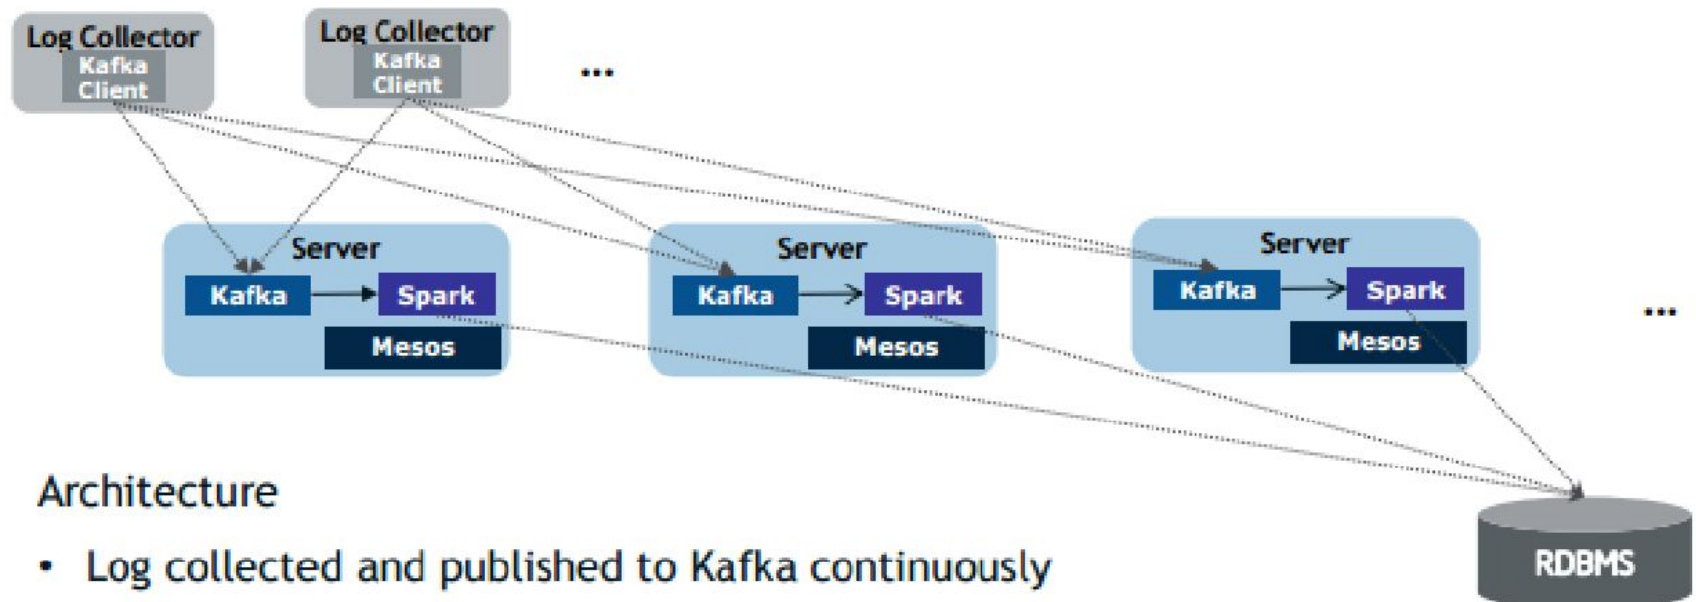


# Real-World Use Case #

## 1 (Semi) Real-Time Log Aggregation & Analysis

- Logs continuously collected & streamed in
- Through queuing/messaging systems Incoming logs processed in a (semi) streaming fashion
- Aggregations for different time periods, demographics, etc.
- Join logs and history tables when necessary  
Aggregation results then consumed in a (semi) streaming fashion
- Monitoring, alerting, etc.

# (Semi) Real-Time Log Aggregation: MiniBatch Jobs




## Architecture

- Log collected and published to Kafka continuously
  - Kafka cluster collocated with Spark Cluster
- Independent Spark applications launched at regular interval
  - A series of mini-batch apps running in “fine-grained” mode on Mesos
  - Process newly received data in Kafka (e.g., aggregation by keys) & write results out

In production in the user's environment today

- 10s of seconds latency

# (Semi) Real-Time Log Aggregation: Spark Streaming



## Implications

- Better streaming framework support – Complex (e.g., stateful) analysis, fault-tolerance, etc.
- Kafka & Spark not collocated – DStream retrieves logs in background (over network) and caches blocks in memory
- Memory tuning to reduce GC is critical – `spark.cleaner.ttl` (throughput \* `spark.cleaner.ttl` < spark mem free size) – Storage level (MEMORY\_ONLY\_SER2)
- Low latency (several seconds) – No startup overhead (reusing SparkContext)